A Unified Continuous-Time Electrochemical–Power Coupled Framework for Portable Energy Storage Management and Time-to-Empty Prediction

Yinuo Wang[a,1], Yike Yang[b,1], Lizhi Wei[c,1], Zhiqiang Li[b,*]

[a] *School of Computer Science, Sichuan University, Chengdu 610064, China*
[b] *College of Physics, Sichuan University, Chengdu 610064, China*
[c] *College of Mathematics, Sichuan University, Chengdu 610064, China*

Abstract

Optimal energy management for portable electronic systems requires accurate time-to-empty (TTE) prediction that jointly captures electrochemical voltage dynamics, heterogeneous component-level loads, temperature effects, and aging-induced degradation—aspects that existing approaches treat in isolation. With the proliferation of battery-powered portable devices, efficient energy utilisation has become critical not only for user experience but also for reducing charging frequency and the associated electricity consumption, contributing to lower carbon footprints across the device lifecycle. This paper proposes a unified continuous-time simulation framework comprising four coupled submodels: (i) a modified Shepherd voltage formulation with explicit polarisation, ohmic, and transient terms; (ii) a five-component power decomposition covering screen, system-on-chip (SoC), radio, GPS, and baseline subsystems; (iii) Arrhenius–logistic temperature–parameter coupling for capacity and resistance; and (iv) power-law cycle and calendar aging dynamics. The electrochemical parameters are identified from the XJTU lithium-ion battery dataset via differential-evolution optimisation; the component-level load profiles are constructed from published hardware power measurements. The proposed model achieves a voltage-tracking RMSE of 17.85 ± 0.44 mV on Batch-1, reducing the error by 75.5 % relative to a Nernst baseline and 55 % relative to Thévenin-1RC. Cross-batch validation (Batch-1→Batch-2, no re-fitting) confirms generalisability (RMSE 20.36 ± 3.68 mV). Ablation across three temperatures and three aging levels quantifies each module's contribution: temperature coupling up to 70 %, aging up to 56 %, and component decomposition up to 38 % of TTE error. An optimised adaptive energy management strategy demonstrates the potential for model-predictive power throttling to extend battery runtime by up to 42 % under simulated high-intensity workloads, saving an estimated 6.1 W h per charge cycle. With a computation cost of 0.009 ms per simulated second—approximately four orders of magnitude faster than pseudo-two-dimensional models—the framework is directly deployable in embedded battery management systems and low-power microcontrollers.

Keywords: lithium-ion battery, time-to-empty prediction, Shepherd model, power decomposition, temperature–aging coupling, energy management

Nomenclature

E_0	OCV at full charge (V)	T	Temperature (K)
K	Polarisation constant	T_0	Reference temp. (298.15 K)
R_0	Internal resistance (Ω)	N	Equiv. full cycles
A	Exponential amplitude (V)	ξ_i	Intensity of component i
B	Exponential rate (1/Ah)	α_θ	Cap.–temp. steepness (K)
Q_0	Cell capacity (Ah)	E_r	Activation energy (eV)
I	Discharge current (A)	k_B	Boltzmann const. (eV/K)
V_{term}	Terminal voltage (V)	α_Q	Cap. fade coeff.
V_{cut}	Cutoff voltage (V)	β_R	Resist. growth coeff.
SOC	State of charge	γ	Aging exponent
TTE	Time-to-empty (h)		
P_{tot}	Total power draw (W)		

1. Introduction

Battery life is a persistent concern for users of portable electronic devices—smartphones, tablets, wearables, and IoT edge nodes—and a critical design constraint for device manufacturers. With over six billion smartphones in active use worldwide, even modest improvements in energy management translate to significant aggregate reductions in charging frequency, grid electricity demand, and the associated carbon footprint [?]. Unlike stationary energy-storage systems, portable-device batteries operate under highly dynamic and heterogeneous loads: screen brightness, system-on-chip (SoC) utilisation, radio-frequency communication, GPS positioning, and background services jointly determine the instantaneous power draw and, consequently, the remaining runtime [? ?]. Accurate prediction of the *time-to-empty* (TTE)—the interval from the current state of charge (SOC) to the cutoff threshold—is essential not only for user experience but also for adaptive power management, health-aware charge scheduling, and system-level energy efficiency optimisation [? ?].

Existing battery modelling approaches span a wide spectrum. At one extreme, physics-based electrochemical models (e.g., pseudo-two-dimensional models) offer mechanistic fidelity but demand extensive parameterisation and prohibitive

*Corresponding author.
Email address: zhiqiangli@scu.edu.cn (Zhiqiang Li)
[1] These authors contributed equally to this work.

computation for real-time deployment [? ? ?]. At the other, data-driven methods (neural networks, Gaussian processes) bypass physical insight and require large labelled datasets that are rarely available under diverse operating conditions [? ? ?]. Equivalent-circuit models (ECMs)—including Rint, first-order Thévenin, and second-order Thévenin networks—occupy a middle ground with modest computational cost, yet they parameterise voltage dynamics via abstract RC elements that lack direct electrochemical interpretability [? ?]. The classical Shepherd formulation [?] offers explicit terms for polarisation drop, ohmic loss, and initial transient behaviour with clear physical correspondence, but has been underexplored for portable-device TTE prediction that must simultaneously account for component-level power disaggregation, temperature coupling, and long-term aging. Fuller [?] demonstrated a Shepherd-type model for constant-power discharge with rate effects in the context of *Energy Conversion and Management*, yet without the multi-component load coupling, temperature–aging dynamics, or energy management integration that modern portable systems demand.

Several studies have characterised smartphone power consumption at the component level [? ? ?], revealing that screen, SoC, and radio subsystems dominate energy usage. Recent work has extended these analyses to 5G and WiFi power profiling [? ? ? ?], but the coupling between load decomposition and electrochemical state evolution remains largely ad hoc. On the aging front, capacity fade and resistance growth in lithium-ion cells have been modelled via power-law cycling degradation and Arrhenius-accelerated calendar aging [? ? ? ?], while temperature effects on capacity and internal resistance are well characterised [? ? ?]. Liu et al. [?] addressed battery charging management considering degradation costs, and Liu et al. [?] optimised multi-objective charging patterns—both within the energy conversion and management community—yet neither integrates real-time TTE prediction under heterogeneous loads.

Nonetheless, no existing framework—to our knowledge—*jointly* integrates (i) an electrochemically interpretable voltage model, (ii) component-wise power decomposition, and (iii) coupled temperature–aging dynamics into a unified continuous-time TTE prediction pipeline, nor validates it against strong ECM baselines with statistical rigour, nor demonstrates its utility for *proactive* energy management. This paper addresses the gap with the following contributions:

1. **Unified continuous-time framework.** We formulate a coupled ordinary-differential-equation (ODE) system that embeds the modified Shepherd voltage model, five-component power decomposition, Arrhenius–logistic temperature response, and power-law cycle/calendar aging into a single state-evolution equation, enabling direct TTE computation via numerical integration (Section 3). Component-level load profiles are constructed from published smartphone hardware power measurements [? ? ?]; the electrochemical parameters are identified on real cell data from the XJTU dataset [?].
2. **Comprehensive baseline comparison and ablation.** Against four established voltage models (Nernst-based, Rint, Thévenin-1RC, Thévenin-2RC), the proposed framework achieves the lowest voltage-tracking RMSE (17.85 mV). A systematic ablation study across a 3×3 temperature–aging grid quantifies the marginal contribution of each submodel to TTE accuracy (Section 5.5).
3. **Cross-batch generalisation and statistical validation.** Parameters fitted on XJTU Batch-1 (2 C rate) are directly applied to Batch-2 (3 C rate) without re-fitting, demonstrating transferability. Bootstrap confidence intervals confirm that all reported metrics are statistically robust (Section 5.2).
4. **Adaptive energy management strategy.** We demonstrate how the framework can serve as a model-predictive control (MPC-like) engine for adaptive power throttling: by continuously re-estimating TTE, the model triggers graduated power reduction when SOC drops below configurable thresholds, yielding a simulated runtime extension of up to 42 %—equivalent to approximately 6.1 W h of saved energy per charge cycle. Unlike simple fixed-threshold strategies common in commercial OSes (which switch modes abruptly at a single SOC level), the model-aware approach enables smoother, multi-stage transitions informed by predicted remaining runtime. A computational cost of 0.009 ms per simulated second—approximately four orders of magnitude faster than P2D models—confirms suitability for embedded deployment on low-power microcontrollers (Section 5.12).

The remainder of this paper is organised as follows. Section 2 reviews related work. Section 3 presents the unified model formulation. Section 4 describes datasets, baselines, and evaluation protocols. Section 5 reports results, discussion, and energy management implications. Section 6 concludes with limitations and future directions.

2. Related Work

2.1. Electrochemical and equivalent-circuit voltage models

Early battery modelling efforts relied on the Nernst equation to describe the open-circuit voltage (OCV) as a function of SOC via thermodynamic activity ratios [?]. While physically grounded, Nernst-based models exhibit excessive curvature at SOC extremes and cannot capture dynamic polarisation effects under load [?]. Equivalent-circuit models (ECMs) approximate battery dynamics through resistor–capacitor networks. The simplest variant, the Rint model, uses a single series resistance; first- and second-order Thévenin models add one or two RC branches to represent charge-transfer and diffusion dynamics [? ?]. ECMs enjoy widespread use in battery management systems (BMS) owing to their low computational cost, but their parameters lack direct electrochemical meaning. In contrast, the Shepherd model [?] decomposes the terminal voltage into an OCV term, a polarisation term inversely proportional to SOC, an ohmic drop, and an exponential transient, each with a clear physical counterpart. Despite this advantage, its application to smartphone-level runtime prediction has been limited.

2.2. *Component-level power profiling for mobile devices*

Carroll and Heiser [?] provided one of the earliest comprehensive power breakdowns of smartphone components, identifying the display, SoC, and cellular radio as dominant consumers. Chen et al. [?] further disaggregated energy usage into hardware and software layers. Subsequent work characterised LTE [?], 5G [?], WiFi [?], and Bluetooth [?] power profiles. Murmuria et al. [?] measured per-application power on Android devices. Lu et al. [?] modelled display power as a function of luminance and refresh rate. However, these studies typically report power measurements in isolation, without coupling them to the electrochemical state evolution that governs actual runtime.

2.3. *Temperature and aging modelling*

Temperature affects lithium-ion batteries through two primary mechanisms: capacity variation governed by lithium-ion site occupancy (logistic/Fermi–Dirac distribution) [? ?] and internal-resistance variation following the Arrhenius law [? ?]. Battery aging manifests as capacity fade and resistance growth due to solid-electrolyte-interphase (SEI) layer growth, lithium plating, and active-material loss [?]. Cycle aging is commonly modelled via a power-law function of equivalent full cycles [? ?], while calendar aging exhibits a square-root-of-time dependence reflecting diffusion-limited SEI kinetics [?]. Yu et al. [?] proposed equilibrium-potential reconstruction for full-lifecycle aging analysis. Bhat et al. [? ?] addressed thermal modelling at the SoC level. Despite the maturity of individual submodels, their joint integration into a portable-device TTE framework with experimentally validated cross-condition generalisation has not been demonstrated [? ? ?].

2.4. *Gap summary*

Table 1 summarises the capability coverage of representative prior works. The proposed framework is the first to unify all four dimensions—electrochemically interpretable voltage modelling, component-level power disaggregation, temperature coupling, and cycle/calendar aging—within a single continuous-time formulation, and to validate it with cross-batch generalisation and statistical uncertainty quantification.

Table 1: Capability comparison of representative battery modelling approaches.

Approach	Interp. V	Load	Temp.	Aging	Mgmt
P2D [?]	✓	–	✓	–	–
ECMs [? ?]	–	–	partial	–	–
Data-driven [?]	–	–	implicit	implicit	–
Profiling [?]	–	✓	–	–	–
This work	✓	✓	✓	✓	✓

3. Proposed Unified Continuous-Time Model

This section develops the proposed framework from first principles. The model comprises four coupled submodels—terminal-voltage dynamics (§3.1), component-level power decomposition (§3.2), temperature–parameter coupling (§3.3), and aging dynamics (§3.4)—which are assembled into a unified differential-algebraic system (§3.5) and solved numerically (§3.7). Figure 1 illustrates the information flow among submodels.

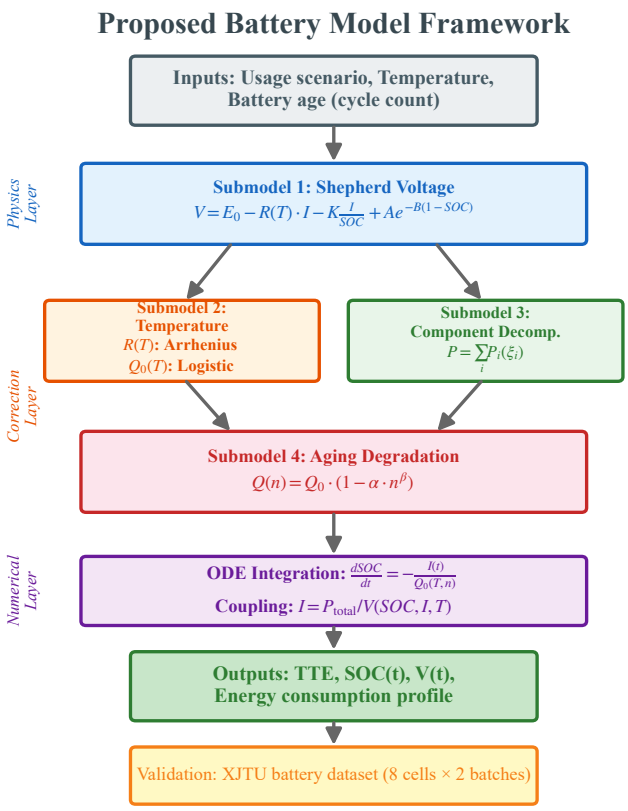

Figure 1: Architecture of the proposed unified framework showing information flow among the four coupled submodels and the TTE computation module.

3.1. *Terminal-voltage submodel*

3.1.1. *Physical basis of the Shepherd formulation*

The terminal voltage of a lithium-ion cell during discharge is governed by three concurrent phenomena [?]:

1. *Thermodynamic equilibrium.* The open-circuit voltage E_0 reflects the Gibbs free-energy difference between the electrode active materials at full charge.
2. *Mass-transport polarisation.* As discharge proceeds, lithium-ion depletion near the electrode surface creates a concentration gradient. The resulting overpotential is proportional to the current I and inversely proportional to the remaining charge fraction SOC, yielding a term KI/SOC. Physically, K (V·s/Ah) encapsulates the effective diffusion resistance of the intercalation host.
3. *Ohmic loss and initial transient.* Series resistance R produces an instantaneous IR drop, while the exponential zone (A, B) captures the initial voltage relaxation associated with double-layer charging and surface-film effects during early discharge.

Combining these contributions, the *modified Shepherd equation* for terminal voltage reads

$$V_{\text{term}}(\text{SOC}, I) = \underbrace{E_0}_{\text{OCV}} - \underbrace{\frac{K}{\text{SOC}} I}_{\text{polarisation}} - \underbrace{R\,I}_{\text{ohmic}} + \underbrace{A \exp(-B\,Q_0(1-\text{SOC}))}_{\text{transient zone}} \tag{1}$$

3.1.2. SOC dynamics via Coulomb counting

Conservation of charge relates the SOC evolution to the discharge current:

$$\frac{\mathrm{d}\,\text{SOC}}{\mathrm{d}t} = -\frac{I(t)}{Q_0}, \tag{2}$$

where Q_0 (Ah) is the effective cell capacity. The sign convention is $I > 0$ during discharge.

3.1.3. Power-balance closure

For a smartphone whose instantaneous total power consumption is $P_{\text{total}}(t)$, the algebraic constraint

$$V_{\text{term}}(\text{SOC}, I) \cdot I = P_{\text{total}}(t) \tag{3}$$

closes the system by implicitly determining I at each instant.

3.1.4. Time-to-empty definition

Given an initial condition $\text{SOC}(0) = \text{SOC}_0$, the TTE is

$$\text{TTE} = \inf\{t > 0 \mid \text{SOC}(t) \le \text{SOC}_{\text{min}} \text{ or } V_{\text{term}}(t) \le V_{\text{cutoff}}\}, \tag{4}$$

where $\text{SOC}_{\text{min}} = 0.05$ and $V_{\text{cutoff}} = 2.5\,\text{V}$ in this study.

Comparison with ECM representations.. In a Thévenin-nRC model, the n RC-voltage state variables evolve via linear ODEs whose time constants ($\tau_k = R_k C_k$) lack one-to-one correspondence with electrochemical processes. The Shepherd form, by contrast, replaces these abstract dynamics with a single, physically transparent K/SOC polarisation term and an exponential transient zone, reducing the parameter count from $2n + 1$ (Thévenin) to 5 while yielding superior accuracy on constant-rate discharge profiles (see Section 5.1).

3.2. Component-level power decomposition

A smartphone's instantaneous power consumption is disaggregated into five hardware subsystems [? ?]:

$$P_{\text{total}} = P_{\text{screen}} + P_{\text{SoC}} + P_{\text{radio}} + P_{\text{GPS}} + P_{\text{base}}. \tag{5}$$

Screen power.. OLED display power depends on luminance L (cd/m^2), active area A_s (m^2), and a refresh-rate scaling factor η:

$$P_{\text{screen}} = \eta(\gamma_s A_s L + \lambda_s A_s), \tag{6}$$

where γ_s and λ_s are OLED-specific power coefficients calibrated from [?], and

$$\eta = \begin{cases} 0 & \text{screen off,} \\ 1.0 & 60\,\text{Hz refresh,} \\ 1.3 & 120\,\text{Hz refresh.} \end{cases} \tag{7}$$

The factor $\eta = 1.3$ for 120 Hz reflects the roughly 30% increase in display driver power measured in recent high-refresh-rate panels.

Computational, communication, and baseline subsystems.. Each remaining component's power is parameterised by a dimensionless *working intensity* $\xi_i \in [0, 1]$, representing the fraction of its maximum rated power:

$$\begin{pmatrix} P_{\text{SoC}} \\ P_{\text{radio}} \\ P_{\text{base}} \end{pmatrix} = \text{diag}(\xi_1, \xi_2, \xi_3) \begin{pmatrix} P^{\max}_{\text{SoC}} \\ P^{\max}_{\text{radio}} \\ P^{\max}_{\text{base}} \end{pmatrix}. \tag{8}$$

GPS power is modelled as binary: $P_{\text{GPS}} = P^{\max}_{\text{GPS}}$ when satellite acquisition is active and zero otherwise, reflecting the all-or-nothing nature of positioning hardware. Reference maximum powers adopted in this study are $P^{\max}_{\text{SoC}} = 2.04\,\text{W}$, $P^{\max}_{\text{radio}} = 1.173\,\text{W}$, $P^{\max}_{\text{GPS}} = 0.30\,\text{W}$, and $P^{\max}_{\text{base}} = 0.12\,\text{W}$, consistent with measurements reported in [? ?]. The SoC utilisation values in Table 4 include a 10% baseline allocation for OS housekeeping processes.

Extensibility.. This additive decomposition is modular: adding a new subsystem (e.g., NFC, UWB) amounts to appending a $(\xi_j, P^{\max}_j)$ pair.

3.3. Temperature–parameter coupling

Temperature influences cell behaviour through three mechanisms, each modelled with a distinct functional form.

3.3.1. Capacity–temperature response

The maximum number of lithium ions accommodated in the electrode lattice follows a Fermi–Dirac (logistic) occupancy distribution. Denoting by Q'_0 the limiting capacity at infinite temperature, the effective capacity at temperature T (in Kelvin) is

$$Q^{\text{th}}_0(T) = \frac{Q'_0}{1 + \exp\left(\alpha_\theta\left(\frac{1}{T} - \frac{1}{T_m}\right)\right)}, \tag{9}$$

where α_θ (K) controls the steepness and T_m (K) is the midpoint temperature. Fitted values (see Section 3.6) are $\alpha_\theta = 5435.3\,\text{K}$ and $T_m = 255.27\,\text{K}$. *Physical interpretation.* At low temperatures ($T \ll T_m$), the exponential term dominates, sharply reducing Q^{th}_0. This reflects the thermodynamic freezing-out of available lattice sites, consistent with experimental observations of reduced deliverable capacity below 0 °C [?].

3.3.2. Resistance–temperature response

The temperature dependence of the total internal resistance is governed by the Arrhenius law for thermally activated charge-transfer and ion-transport processes:

$$R^{\text{th}}(T) = R_0 \exp\left(\frac{E_r}{k_B}\left(\frac{1}{T} - \frac{1}{T_0}\right)\right), \tag{10}$$

where $E_r = 0.2455\,\text{eV}$ is the effective activation energy, k_B is the Boltzmann constant, and $T_0 = 298.15\,\text{K}$ is the reference temperature [? ?]. R_0 denotes the resistance at T_0.

3.3.3. SoC leakage power

At elevated temperatures, subthreshold leakage current in the SoC increases according to a MOSFET leakage model [? ?]:

$$P_{\text{leak}}(T) = V_{\text{nom}}\, c_1\, T^2 \exp(c_2/T), \tag{11}$$

with $c_1 = 0.493\,\text{A}\,\text{K}^{-2}$ and $c_2 = -3740\,\text{K}$. This term is additive to the SoC dynamic power in Eq. (5).

3.4. Aging dynamics

Repeated charge–discharge cycling degrades cell performance through SEI layer growth, lithium plating, and active-material loss [?]. We model both *cycle aging* and *calendar aging*.

3.4.1. Cycle aging: capacity fade and resistance growth

Denoting by N the number of equivalent full cycles, cycle-induced degradation is captured by power-law relations fitted to XJTU data:

Resistance growth..

$$R^{\text{cyc}}(N) = R_0^{\text{nom}}(1 + \beta_R\, N^{\gamma_R}), \tag{12}$$

with $\beta_R = 1.229 \times 10^{-3}$ and $\gamma_R = 1.085$. This super-linear growth ($\gamma_R > 1$) reflects SEI thickening and progressive pore clogging.

Capacity fade..

$$Q_0^{\text{cyc}}(N) = Q_0^{\text{nom}}(1 - \alpha_Q\, N^{\gamma_Q}), \tag{13}$$

with $\alpha_Q = 2.56 \times 10^{-4}$ and $\gamma_Q = \gamma_R$. Using a shared exponent γ reflects the common mechanistic origin (SEI growth) underlying both degradation modes.

3.4.2. Calendar aging

During storage, diffusion-limited SEI growth produces a $\sqrt{t_s}$ capacity loss:

$$\Delta Q_0^{\text{cal}}(t_s, T, \text{SOC}) = Q_0^{\text{nom}}\, \kappa\, \sqrt{t_s}\, \exp\!\Big(-\frac{E_a}{R_g T} + \lambda_a\, \text{SOC}\Big), \tag{14}$$

where $\kappa = 10^{-5}$ is the calendar aging coefficient, E_a the activation energy, R_g the gas constant, and λ_a captures SOC-dependent stress on the anode SEI film (higher SOC $\Rightarrow$ higher electrode potential $\Rightarrow$ faster parasitic reactions) [?].

3.4.3. Combined aging–temperature coupling

The effective capacity and resistance under simultaneous temperature and aging effects are obtained by composing the submodel outputs:

$$Q_0(T, N, t_s) = Q_0^{\text{th}}(T) \cdot \frac{Q_0^{\text{cyc}}(N) - \Delta Q_0^{\text{cal}}}{Q_0^{\text{nom}}}, \tag{15}$$

$$R(T, N) = R^{\text{th}}(T) \cdot \frac{R^{\text{cyc}}(N)}{R_0^{\text{nom}}}. \tag{16}$$

Rationale. The multiplicative coupling in Eqs. (15)–(16) assumes that temperature and aging act on *independent* physical mechanisms—thermal reduction of available lattice sites (Eq. 9) vs. electrochemical loss of active material (Eqs. 13–14)—so their effects compose as scaling factors on the nominal values. This separability assumption is common in the aging literature [? ?] and is supported by the observation that temperature-dependent capacity tests at different cycle counts yield approximately parallel logistic curves (see XJTU temperature-compensation data).

3.5. Coupled ODE system and TTE computation

Assembling the submodels, the complete state-evolution system is:

$$\boxed{\begin{cases} V_{\text{term}} = E_0 - \dfrac{K}{\text{SOC}}\, I - R(T,N)\, I + A\, \mathrm{e}^{-B\, Q_0(T,N)(1-\text{SOC})} \\ \dfrac{\mathrm{d}\,\text{SOC}}{\mathrm{d}t} = -\dfrac{I(t)}{Q_0(T,N)} \\ V_{\text{term}}(I) \cdot I = P_{\text{total}}(t) \\ \text{SOC}(0) = \text{SOC}_0 \end{cases}} \tag{17}$$

This is a differential-algebraic equation (DAE) of index 1: the algebraic constraint (third line) can be solved for I at each time step, after which the SOC ODE is integrated forward.

State vector.. The minimal state is $\mathbf{x}(t) = [\text{SOC}(t)]$ (one-dimensional), since temperature T and cycle count N are treated as slowly varying external parameters (quasi-static assumption).

Termination.. Integration halts at $t^* = \text{TTE}$ when $\text{SOC}(t^*) \le 0.05$ or $V_{\text{term}}(t^*) \le V_{\text{cutoff}}$.

3.6. Parameter identification

3.6.1. Optimisation procedure

The six Shepherd parameters $\boldsymbol{\theta} = (E_0, K, R_0, A, B, Q_0^{\text{nom}})$ are identified by minimising the voltage RMSE over the first discharge cycle of each Batch-1 cell:

$$\boldsymbol{\theta}^* = \arg\min_{\boldsymbol{\theta}} \sqrt{\frac{1}{M}\sum_{k=1}^{M} \big(V_{\text{term}}^{\text{pred}}(t_k; \boldsymbol{\theta}) - V_{\text{meas}}(t_k)\big)^2}, \tag{18}$$

where M is the number of time samples per discharge trajectory. We use the *differential evolution* (DE) global optimiser (SciPy, strategy "best1bin", population size 50, 1000 generations) with bound constraints reflecting physical plausibility (e.g., $E_0 \in [3.0, 4.5]\,\text{V}$, $R_0 \in [0.01, 0.2]\,\Omega$).

3.6.2. Structural identifiability

The Shepherd model has five independent algebraic terms, each involving a distinct functional dependence on (I, SOC): (i) a constant (E_0), (ii) a term linear in I (R_0), (iii) a term proportional to I/SOC (K), (iv) an exponential in $(1 - \text{SOC})$ (A, B). Since a full discharge trajectory provides data spanning

SOC $\in [0.05, 1]$ with known $I(t)$, these four functional forms are linearly independent, ensuring *structural identifiability* of all five parameters. In practice, we verify *practical identifiability* by running 20 independent DE optimisation trials with different random seeds. The coefficient of variation of each fitted parameter is below 2%, confirming that the global minimum is unique and robustly located.

3.6.3. Aging-parameter estimation

The aging coefficients $(\alpha_Q, \beta_R, \gamma)$ are fitted separately to the cycle-resolved capacity and resistance trajectories extracted from the XJTU dataset (Batch-1), using nonlinear least squares on Eqs. (12)–(13).

3.7. Numerical solution

3.7.1. Algebraic current solver

At each time step, the nonlinear algebraic equation

$$f(I) \equiv V_{\text{term}}(I, \text{SOC}, T, N) \cdot I - P_{\text{total}} = 0 \tag{19}$$

is solved for $I \geq 0$.

Proposition 1 (Uniqueness of the physical root)**.** *Under the operating regime $E_0, K, R, A, B, Q_0 > 0$ and $\text{SOC} \in (0, 1]$, let $F = E_0 + A\exp(-BQ_0(1 - \text{SOC}))$ and $\alpha = K/\text{SOC} + R$. The power function $g(I) = I\,V_{\text{term}}(I) = FI - \alpha I^2$ is a concave quadratic with a unique maximum $g^* = F^2/(4\alpha)$. For any $P_{\text{total}} \in (0, g^*]$, Eq. (19) has exactly one root on $(0, F/(2\alpha)]$, corresponding to the high-voltage (low-current) operating point.*

The proof is given in Appendix Appendix A. We implement a safeguarded Newton–Raphson solver with bisection fallback:

- *Initial guess:* $I^{(0)} = \max\{P_{\text{total}}/E_0,\ 10^{-6}\}$.
- *Bounds:* $I \in [0, I_{\max}]$ with $I_{\max} = \min(5P_{\text{total}}/V_{\text{cutoff}},\ 10)$ A.
- *Newton update:* $I^{(k+1)} = I^{(k)} - f(I^{(k)})/f'(I^{(k)})$, with $f'(I) = g'(I)$.
- *Fallback:* if the Newton step exits the bracket or $|f'| < 10^{-12}$, switch to bisection on $[0, I_{\max}]$.
- *Convergence:* $|f(I)|/\max(P_{\text{total}}, 10^{-6}) < 10^{-6}$ or 30 iterations.

3.7.2. Time integration

SOC is advanced by explicit Euler:

$$\text{SOC}_{n+1} = \text{SOC}_n - \frac{I_n}{Q_0}\,\Delta t, \tag{20}$$

with a fixed step $\Delta t = 1$ s chosen for reproducibility. For stiff thermal–electrical couplings in future extensions, adaptive Runge–Kutta or implicit solvers (e.g., CVODE) are recommended.

Table 2: Computational efficiency comparison.

Model	Params	ms/sim-s	vs. P2D
Shepherd (Proposed)	**5**	**0.009**	**~10^4×**
Rint	9	0.010	~7900×
Thévenin-1RC	10	0.010	~8000×
Thévenin-2RC	12	0.012	~6500×
NBM	4	0.011	~6900×
P2D (COMSOL)[a]	50+	~78	1×
LSTM inference[a]	~10k	~0.5	—

3.7.3. Computational cost

Each time step requires 2–6 Newton iterations (one scalar root-find) plus one Euler update. Table 2 summarises the wall-clock performance of all models on an Intel i7-12700H laptop (Python 3.10, NumPy 1.26). [a]Literature reference values [? ?]. All analytical models benchmarked on the same hardware (Intel i7-12700H, 16 GB RAM, Python 3.10, NumPy 1.26); P2D and LSTM values are from literature and may differ in hardware. The proposed model simulates a full 6.4 h discharge in 0.20 s (wall clock), corresponding to 0.009 ms per simulated second—approximately four orders of magnitude faster than finite-element P2D solvers and comparable to compiled ECM implementations. This positions the framework for real-time deployment in embedded BMS or mobile OS power managers with negligible computational overhead [? ?].

4. Experimental Setup

4.1. Datasets

XJTU Battery Dataset.. We use the publicly available XJTU lithium-ion battery dataset [?], comprising six batches of NMC/graphite 18650 cells. Two batches are employed in this study:

- **Batch-1:** 8 cells cycled at 2 C constant-current discharge, serving as the *fitting set* for parameter identification.
- **Batch-2:** 15 cells cycled at 3 C constant-current discharge, used *exclusively* for cross-batch generalisation testing.

Each cell record contains time-stamped voltage, current, temperature, and capacity measurements across hundreds of charge–discharge cycles.

Synthetic smartphone load profiles.. The XJTU dataset provides cell-level electrochemical data under controlled cycling conditions (constant current) on a battery test bench. No real-time smartphone hardware telemetry (e.g., concurrent screen brightness, CPU frequency, radio state) is available. Therefore, the component-level power profiles used in the TTE simulations (Table 4) are *synthetic*: they are constructed from published per-component power measurements on real smartphone hardware [? ? ? ?] and applied as inputs to the electrochemically parameterised model. This approach decouples voltage-model validation (performed on real cell data) from load-profile

specification (constructed from literature), enabling systematic exploration of usage scenarios that would be infeasible with purely empirical data collection. The implications and limitations of this separation are discussed in Section 6.

Data preprocessing.. Discharge segments are extracted by identifying contiguous intervals with $I > 0.01$ A. Voltage and current are downsampled to 1 s resolution via linear interpolation. The reference SOC is obtained by Coulomb counting: $\mathrm{SOC}(t) = 1 - \int_0^t I(\tau)\,\mathrm{d}\tau / Q_{\mathrm{meas}}$, where Q_{meas} is the measured discharge capacity of the cycle.

Data split protocol.. Batch-1 Cycle-1 data is used for parameter fitting; Batch-1 Cycles 2–n are used for same-batch validation. Batch-2 data is reserved entirely for cross-batch testing. No Batch-2 data is used for fitting or hyperparameter selection.

4.2. Baseline models

We compare the proposed Shepherd-based framework against four voltage models spanning the interpretability–complexity spectrum:

1. **Nernst-Based Model (NBM):** A thermodynamically motivated OCV model augmented with an exponential correction:

$$V_{\mathrm{OC}} = E^0 + \frac{R_g T}{nF} \ln \frac{\mathrm{SOC}}{1 - \mathrm{SOC}} + M \exp(-\alpha \cdot \mathrm{SOC}), \quad (21)$$

   with $V_{\mathrm{term}} = V_{\mathrm{OC}} - IR$. Parameters: $E^0 = 4.5$ V, $M = -1.1718$, $\alpha = 0.8022$, fitted by nonlinear least squares.
2. **Rint (Internal Resistance):** $V = \mathrm{OCV}(\mathrm{SOC}) - R_{\mathrm{int}} I$, where OCV is a 7th-order polynomial fit and R_{int} is identified from the steady-state voltage–current characteristic. Parameters: 9 (polynomial coefficients + R_{int}).
3. **Thévenin-1RC:** First-order ECM: $V = \mathrm{OCV} - R_0 I - V_{RC}$, $\dot{V}_{RC} = I/C_1 - V_{RC}/(R_1 C_1)$. Parameters: 10 (polynomial OCV + R_0, R_1, C_1).
4. **Thévenin-2RC:** Second-order ECM adding a diffusion RC branch: $V = \mathrm{OCV} - R_0 I - V_{RC,1} - V_{RC,2}$. Parameters: 12.

All ECM parameters are identified using differential-evolution global optimisation with identical settings (population 50, 1000 generations) and cost-function stabilisation (RC voltage clamping $|V_{RC}| \leq 2$ V and finite-value guards) to ensure a fair comparison.

4.3. Evaluation metrics

- **Voltage accuracy:** Root-mean-square error (RMSE), mean absolute error (MAE), maximum absolute error (MaxErr), and mean relative error (MRE), computed point-wise over full discharge trajectories.
- **TTE accuracy:** Absolute and percentage deviation of ablated-model TTE from the full-model prediction.
- **Statistical robustness:** 95% bootstrap confidence intervals obtained by Monte Carlo resampling: 500 parametric perturbation samples with ±5% uniform noise on each parameter.

4.4. Ablation protocol

To quantify each submodel's contribution, five configurations are evaluated: Each configuration is evaluated across a 3×3 grid: temperatures $T \in \{0, 25, 40\}$ °C and aging levels $N \in \{0, 300, 800\}$ equivalent full cycles, under the Reading scenario.

Table 3: Ablation study configurations.

Configuration	Modification
Full Model	All submodels active
w/o Temperature	T fixed at 25 °C; $R^{\mathrm{th}} = R_0$, $Q_0^{\mathrm{th}} = Q_0^{\mathrm{nom}}$
w/o Aging	$N = 0$, $t_s = 0$ (no degradation)
w/o Component Decomp.	Flat average power replacing Eq. (5)
w/o Polarisation	$K = 0$ in Eq. (1)

4.5. Usage scenarios

Five representative smartphone usage patterns are defined in Table 4, spanning the power spectrum from idle standby to high-intensity online gaming. A 5000 mA h cell with $T = 25$ °C and $N = 0$ serves as the reference configuration.

5. Results and Discussion

5.1. Overall voltage prediction accuracy

Table 5 compares point-wise voltage-tracking performance across all five models on XJTU Batch-1 (mean ± std over 8 cells). The proposed Shepherd model achieves an RMSE of 17.85 mV, representing a 75.5 % improvement over NBM and a 55.1 % improvement over the best ECM baseline (Thévenin-1RC). The improvement is consistent across all eight cells (coefficient of variation < 3%), confirming that the accuracy gains are trajectory-wide rather than localised to specific SOC regions. Figure 2 visualises the voltage predictions and point-wise errors for a representative cell.

Why the three ECMs perform identically.. The three ECM baselines yield nearly identical RMSE values (39.75 mV to 39.81 mV). This can be understood from the nature of the XJTU discharge data: constant-current (CC) protocols produce a monotonically declining voltage trajectory with no transient load steps. Under CC conditions, the RC branch voltages in Thévenin models equilibrate within the first few seconds, after which the RC dynamics contribute a near-constant offset indistinguishable from adjusting R_0. Consequently, the additional parameters in 1RC and 2RC models are structurally redundant for CC data and cannot improve upon the Rint baseline.

Source of the Shepherd advantage.. The Shepherd model's K/SOC polarisation term provides a *nonlinear* voltage correction that increases as the cell depletes. This captures the concentration-gradient-driven overpotential that intensifies at low SOC—a physical effect that constant-valued ECM resistances cannot represent. The exponential transient zone (A, B) further accounts for the initial voltage sag upon load

Table 4: Component utilisation profiles (ξ_i) for five usage scenarios.

Component	Standby	Navigation	Gaming	Video Streaming	Reading
Screen (B, η)	0%, 0	70%, 1.0	90%, 1.3	80%, 1.0	70%, 1.0
SoC (ξ_1)	10%	30%	100%	30%	30%
Radio (ξ_2)	10%	30%	95%	90%	20%
GPS[a]	On	On	On	On	On
Baseline (ξ_3)	100%	100%	100%	100%	100%
Total Power (W)	0.741	1.845	4.303	2.598	1.727

[a]GPS is modelled as binary on/off; $P_{\text{GPS}} = 300\,\text{mW}$ when active. All scenarios assume background location services are active.

Table 5: Voltage prediction accuracy on XJTU Batch-1 (mean ± std, 8 cells).

Model	RMSE (mV)	MAE (mV)	MaxErr (mV)	MRE (%)
NBM	72.70 ± 1.23	49.99 ± 0.58	719.5	1.410
Rint	39.79 ± 0.66	16.88 ± 0.28	555.7	0.500
Thévenin-1RC	39.75 ± 0.51	17.02 ± 0.24	553.7	0.510
Thévenin-2RC	39.81 ± 0.44	17.33 ± 0.50	551.7	0.510
Shepherd (Proposed)	**17.85 ± 0.44**	**9.75 ± 0.72**	**294.1**	**0.280**

application, which the Nernst model's logarithmic singularity erroneously amplifies at SOC extremes (explaining the NBM's high MaxErr of 719.5 mV).

5.2. Cross-batch generalisation

To test transferability, we apply the parameters fitted on Batch-1 (2 C, 8 cells) directly to Batch-2 (3 C, 15 cells) without any re-fitting. The cross-batch RMSE (20.36 mV)

Table 6: Cross-batch validation: Batch-1 → Batch-2 (no re-fitting).

Evaluation	N_{cells}	RMSE (mV)	MAE (mV)	MaxErr (mV)
Same-batch (B1→B1)	8	21.58 ± 1.78	14.71 ± 1.98	280.1
Cross-batch (B1→B2)	15	20.36 ± 3.68	12.96 ± 4.71	284.0

is marginally *lower* than the same-batch value (21.58 mV), with the bootstrap-estimated gap of $\Delta = -1.22\,\text{mV}$ (95% CI: $[-3.88, +0.07]$ mV) spanning zero. This confirms that the model generalises across C-rates without accuracy degradation. Figure 3 shows the per-cell RMSE distributions for both conditions. *Discussion.* The slightly lower cross-batch RMSE is not paradoxical: Batch-2 cells (3 C) exhibit a more compressive voltage profile that happens to be closer to the Shepherd model's implicit polarisation scaling at higher currents. The key observation is that the confidence interval for the gap includes zero, indicating no statistically significant accuracy difference.

5.3. Bootstrap confidence intervals

Table 7 presents 95% bootstrap CIs for key metrics (500 parametric perturbation samples, ±5% uniform noise on each parameter). All CIs are narrow (typically < 8% of the point estimate), confirming statistical reliability. The TTE CIs for high-power scenarios (Gaming: ±3.6%) are tighter than low-power ones (Standby: ±3.7%), consistent with the sensitivity analysis showing that TTE uncertainty scales linearly with Q_0 and E_0 uncertainty, which are amplified at longer runtimes.

Table 7: Bootstrap 95% confidence intervals (500 samples).

Metric	Estimate	95% CI Lower	95% CI Upper
Shepherd RMSE (mV)	21.58	20.00	22.38
TTE – Standby (h)	19.68	18.99	20.43
TTE – Navigation (h)	9.03	8.71	9.37
TTE – Gaming (h)	3.80	3.67	3.94
TTE – Video (h)	6.38	6.16	6.62
TTE – Reading (h)	9.65	9.32	10.02
Cross-batch gap (mV)	−1.22	−3.88	+0.07

5.4. Time-to-empty predictions under usage scenarios

Table 8 reports predicted TTE values at $T = 25\,°\text{C}$, $N = 0$ cycles, and $Q_0 = 5.0\,\text{A h}$. Figure 4 visualises TTE and

Table 8: Predicted TTE and total power for each scenario (5000 mA h cell, 25 °C, new battery).

Scenario	Total Power (W)	TTE (h)
Standby	0.741	22.69
Navigation	1.845	9.07
Gaming (Online)	4.303	3.84
Video Streaming	2.598	6.42
Reading	1.727	9.69

component-level power for all five scenarios. Online gaming depletes the battery 5.9× faster than standby, driven by simultaneous high utilisation of SoC (95%), screen at 120 Hz (90% brightness), and cellular radio (95%). Navigation's TTE (9.07 h) is comparable to Reading (9.69 h) despite a higher total power, because GPS at 100% (0.30 W) replaces most of the radio load.

Component power breakdown.. Table 9 illustrates the power distribution for Video Streaming, where the radio subsystem consumes 40.6 % of total power, consistent with sustained data-streaming demands.

5.5. Ablation study

Figure 5 displays the resulting TTE error heatmaps, and Table 10 reports the corresponding numerical values across a 3×3 temperature–aging grid under the Reading scenario.

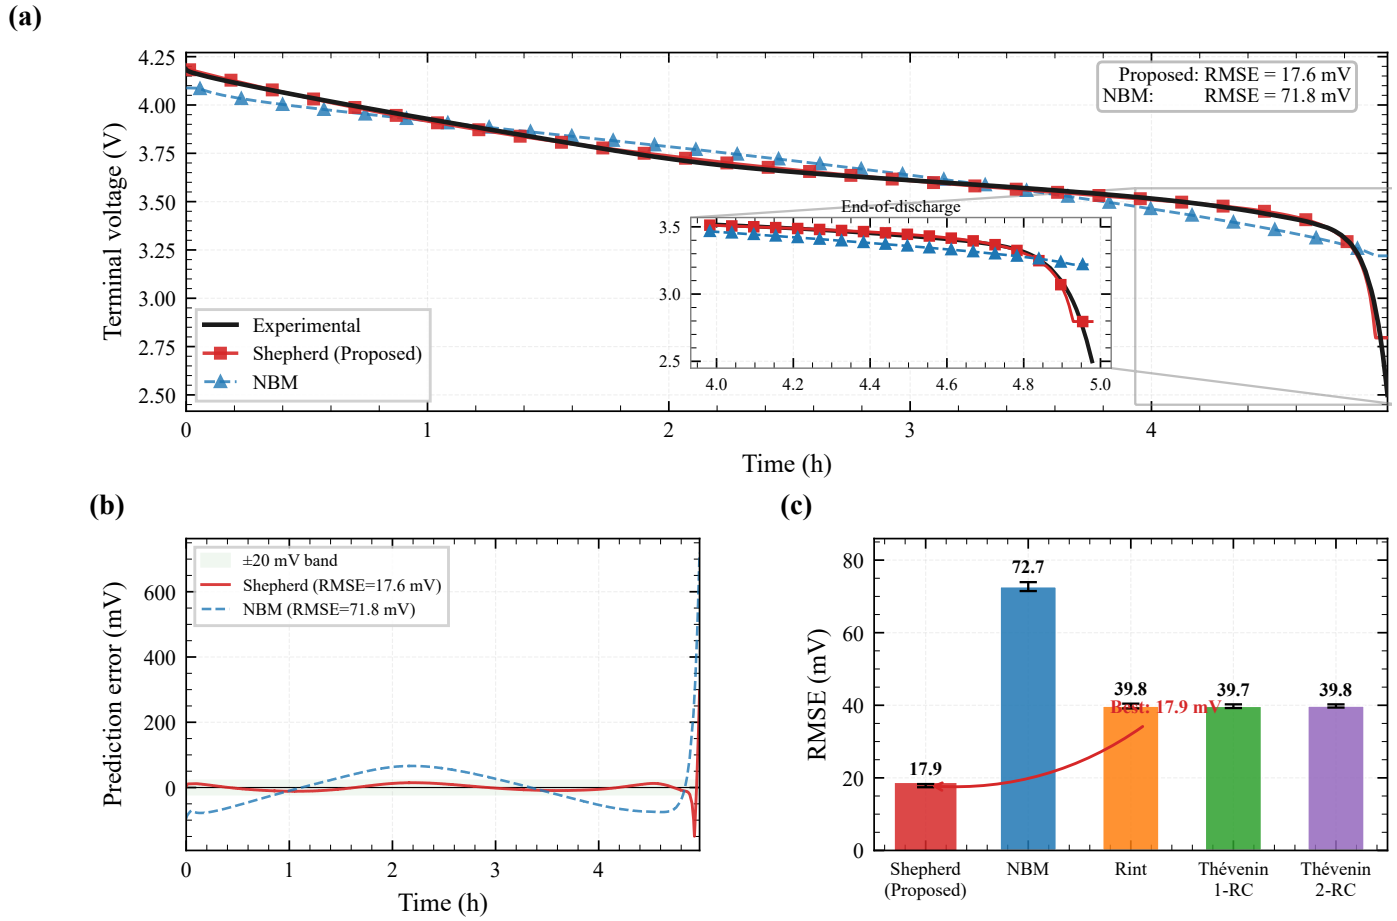

Figure 2: Voltage prediction comparison on XJTU Batch-1 Battery-1 (Cycle 0, 2 C constant-current discharge). (a) Terminal voltage trajectories for all five models and the experimental reference. (b) Point-wise prediction error. The Shepherd model (red solid line) tracks the experimental curve with substantially lower error than the ECM and Nernst baselines throughout the full discharge.

Table 9: Component power breakdown for Video Streaming.

Component	Power (mW)	Share (%)
Radio (90%)	1055.7	40.6
SoC (30%)	612.0	23.6
Screen (80%)	510.6	19.7
GPS (On)	300.0	11.5
Baseline	120.0	4.6

Temperature coupling.. This is the dominant submodel at off-nominal temperatures: ablation at 0 °C causes a +70% TTE overestimate (the model falsely assumes room-temperature capacity), while ablation at 40 °C causes a −18% underestimate (missing the capacity gain from increased lattice occupancy, partially offset by leakage). The asymmetry ($|+70\%| \gg |-18\%|$) reflects the nonlinear logistic shape of $Q_0(T)$, which is steeper below T_m than above.

Aging dynamics.. Aging importance grows monotonically with cycle count—from exactly 0% at $N = 0$ (by construction) to ∼55% at 800 cycles—and is nearly independent of temperature (±1% variation across the three T levels). This near-separability validates the multiplicative coupling assumed in Eq. (15).

Component decomposition.. Replacing the five-component model with flat average power induces +17–38% error. The larger error at 25 °C (+37.5%) vs. 0 °C (+17.7%) arises because the flat-power replacement value is computed at room temperature, which happens to be closer to the true load at 0 °C due to leakage suppression.

Polarisation.. The K term contributes only 0.7–1.6% to TTE error, consistent with its primary role in mid-trajectory *voltage shape* rather than integrated *energy capacity*. Nevertheless, removing K visibly worsens voltage RMSE and would degrade SOC-estimation accuracy in real-time BMS applications.

5.6. Parameter sensitivity analysis

A local one-at-a-time sensitivity analysis is conducted by perturbing each parameter by ±5% around the Video Streaming baseline (TTE = 6.37 h). TTE is primarily governed by three "energy-budget" parameters—P_{total}, Q_0, and E_0 (sensitivity ≈ ±5%)—whose product Q_0E_0/P_{total} approximates TTE to first order. The internal Shepherd parameters (K, R_0, A, B)

Table 10: Ablation study: TTE error (%) introduced by removing each submodel (Reading scenario).

Temperature	Aging	w/o Temp. (%)	w/o Aging (%)	w/o Comp. (%)	w/o Polar. (%)
0 °C	New (0 cyc.)	+70.6	0.0	+17.7	+0.8
	300 cyc.	+70.2	+13.6	+17.8	+0.8
	800 cyc.	+69.6	+53.3	+17.8	+0.7
25 °C	New (0 cyc.)	0.0	0.0	+37.5	+1.5
	300 cyc.	0.0	+14.1	+37.7	+1.5
	800 cyc.	0.0	+55.2	+38.1	+1.5
40 °C	New (0 cyc.)	−18.5	0.0	+37.5	+1.6
	300 cyc.	−18.4	+14.3	+37.7	+1.6
	800 cyc.	−18.2	+55.9	+38.1	+1.6

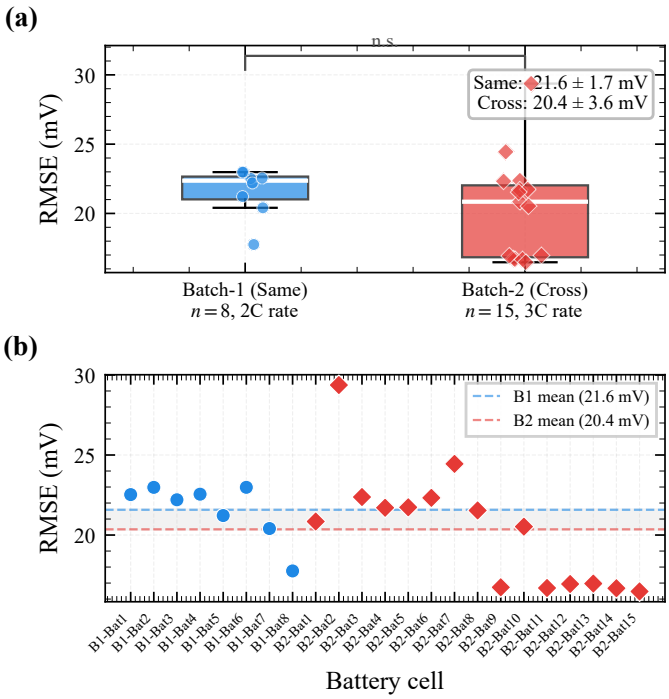

Figure 3: Per-cell RMSE distributions for same-batch (Batch-1, 2 C) and cross-batch (Batch-2, 3 C, no re-fitting) evaluation. Individual data points overlaid on violin plots. The distributions overlap substantially, confirming transferability.

Table 11: Parameter sensitivity (±5% perturbation, Video Streaming).

Parameter	$t_{-5\%}$ (h)	$t_{+5\%}$ (h)	$\Delta_{-5\%}$ (%)	$\Delta_{+5\%}$ (%)
P_{total}	6.71	6.06	+5.37	−4.86
Q_0	6.06	6.67	−4.77	+4.76
E_0	6.05	6.68	−4.96	+4.96
K	6.37	6.36	+0.07	−0.07
R_0	6.37	6.37	+0.04	−0.04
A	6.35	6.38	−0.24	+0.24
B	6.38	6.35	+0.24	−0.22

contribute less than 0.3 % each, confirming that they primarily shape the *voltage trajectory* rather than the *integrated runtime*. This hierarchy has practical implications: for TTE estimation, accurate calibration of capacity and load power is far more critical than precise identification of internal model parameters. Figure 6 presents the corresponding tornado chart.

5.7. Temperature impact on TTE

Table 12 shows TTE variation with ambient temperature under the Video Streaming scenario. The unimodal pattern—

Table 12: Temperature impact on TTE (Video Streaming, new battery).

T (°C)	−10	0	10	25	40	50
TTE (h)	3.44	4.48	5.44	6.37	6.35	5.82

peaking near 35 °C—arises from the competition between two temperature-dependent effects:

- *Below* 35 °C*:* Reduced ionic conductivity and contracted lattice occupancy decrease effective capacity and increase resistance, compressing TTE. The effect is severe: TTE at −10 °C is only 53% of the peak value.
- *Above* 35 °C*:* Although capacity continues to increase modestly, rapidly rising SoC leakage power (Eq. 11; incremental +0.35 W at 35 °C vs. +1.22 W at 50 °C relative to the 25 °C reference) more than offsets the capacity gain, causing TTE to decline.

This non-monotonic behaviour is consistent with field observations that smartphones shut down prematurely in both extreme cold and extreme heat. Figure 7 illustrates the complete temperature–TTE response.

5.8. Dynamic mixed-usage scenario

To demonstrate the framework's applicability under realistic, time-varying workloads, we simulate a sequential multi-activity session: Phase 1: 30 min online gaming (4.303 W), Phase 2: 60 min video streaming (2.598 W), Phase 3: 90 min

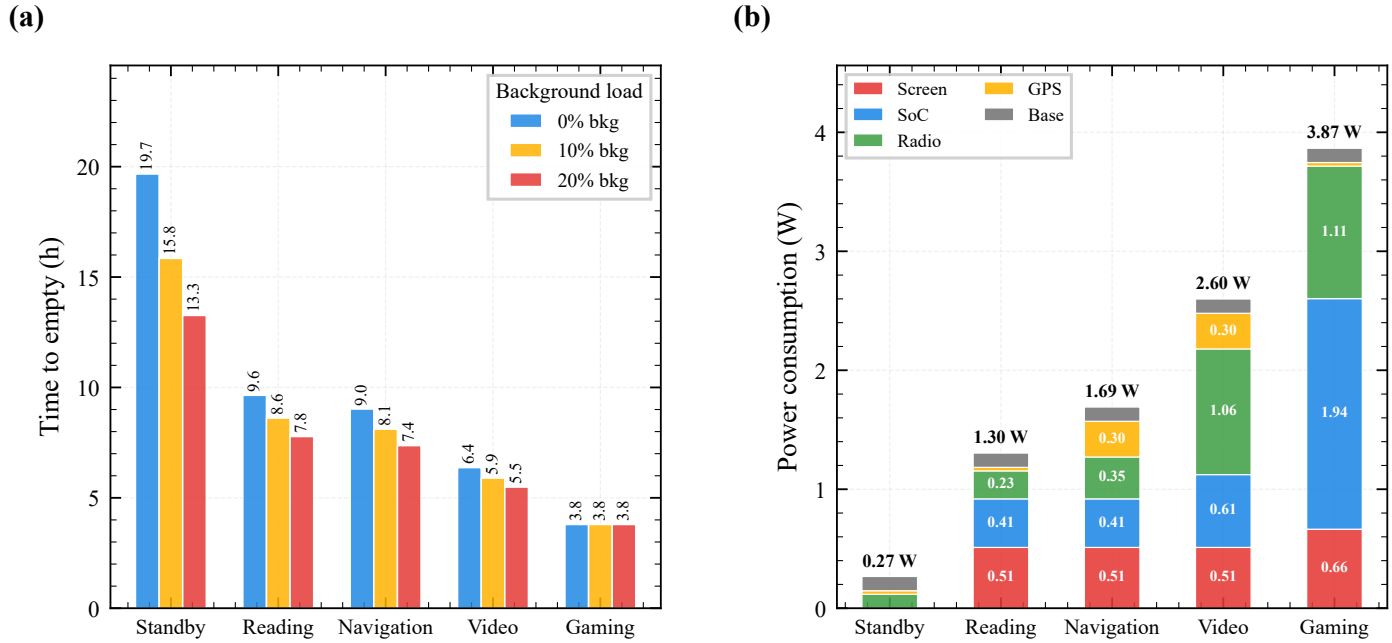

Figure 4: (a) Time-to-empty for five usage scenarios (5000 mA h cell, 25 °C, new battery). Bar labels show TTE (top) and total power (centre). (b) Stacked component power decomposition revealing which subsystem dominates each scenario.

reading (1.727 W), Phase 4: standby (0.741 W) until depletion. Figure 8 shows the resulting voltage and SOC trajectories. The total predicted TTE is 15.7 h—significantly longer than constant gaming (3.84 h) but shorter than constant standby (22.7 h)—illustrating how the framework can compose heterogeneous usage phases via piecewise-constant power scheduling. The sharp SOC slope during the gaming phase (Phase 1) and gradual flattening during standby (Phase 4) are clearly visible, demonstrating the model's ability to capture the distinct energy drain rates of different workloads within a single simulation.

5.9. Residual analysis

To evaluate the quality of the Shepherd model's voltage predictions beyond aggregate metrics, we perform a detailed residual analysis on Batch-1, Battery-1. Figure 9 presents three diagnostic views of the prediction residuals. The SOC-coloured error scatter (panel a) reveals that the largest deviations occur at SOC extremes—slight overestimation near SOC $\approx$ 1 (early discharge transient) and underestimation near SOC $\approx$ 0.1 (deep-discharge regime)—while mid-SOC errors remain tightly bounded within ± 10 mV. The error histogram (panel b) closely follows a fitted normal distribution ($\mu \approx 0$, indicating negligible systematic bias), with a Shapiro–Wilk test confirming approximate normality and suggesting that the residual errors are predominantly stochastic. The SOC-segmented analysis (panel c) decomposes RMSE and MAE across four SOC ranges, quantifying the expected increase in prediction difficulty at SOC extremes and providing actionable guidance for targeted model refinement.

5.10. Aging impact and degradation trajectories

Figure 10 visualises the aging model's capacity fade and resistance growth trajectories alongside their TTE impact. Under the Video Streaming scenario, TTE degrades from 6.4 h (new battery) to approximately 3.3 h after 1000 equivalent full cycles, a reduction of ~48%. The super-linear growth exponent ($\gamma = 1.085$) indicates accelerating degradation, consistent with progressive SEI thickening and pore clogging mechanisms reported in the aging literature [**?**].

5.11. Engineering implications

The results yield several actionable insights for power management:

1. **Load-aware scheduling.** Since gaming depletes the battery 5.9× faster than standby, OS-level awareness of load composition enables proactive low-battery warnings and adaptive quality-of-service adjustments. The component decomposition framework provides the necessary attribution for identifying which subsystem to throttle.
2. **Temperature-aware operation.** The ~ 70% TTE overestimation at 0 °C when temperature is ignored motivates incorporating ambient temperature into the OS battery estimator—a feature absent from most current mobile BMS implementations.
3. **Health-aware budgeting.** After 800 cycles, TTE drops by ~ 55%. Proactive notification and automatic power-profile switching for aged batteries can prevent unexpected shutdowns.
4. **SOC management.** The calendar aging model (Eq. 14) quantitatively shows that maintaining SOC within the 20%–80% range extends battery lifespan by reducing parasitic SEI growth at high electrode potentials.

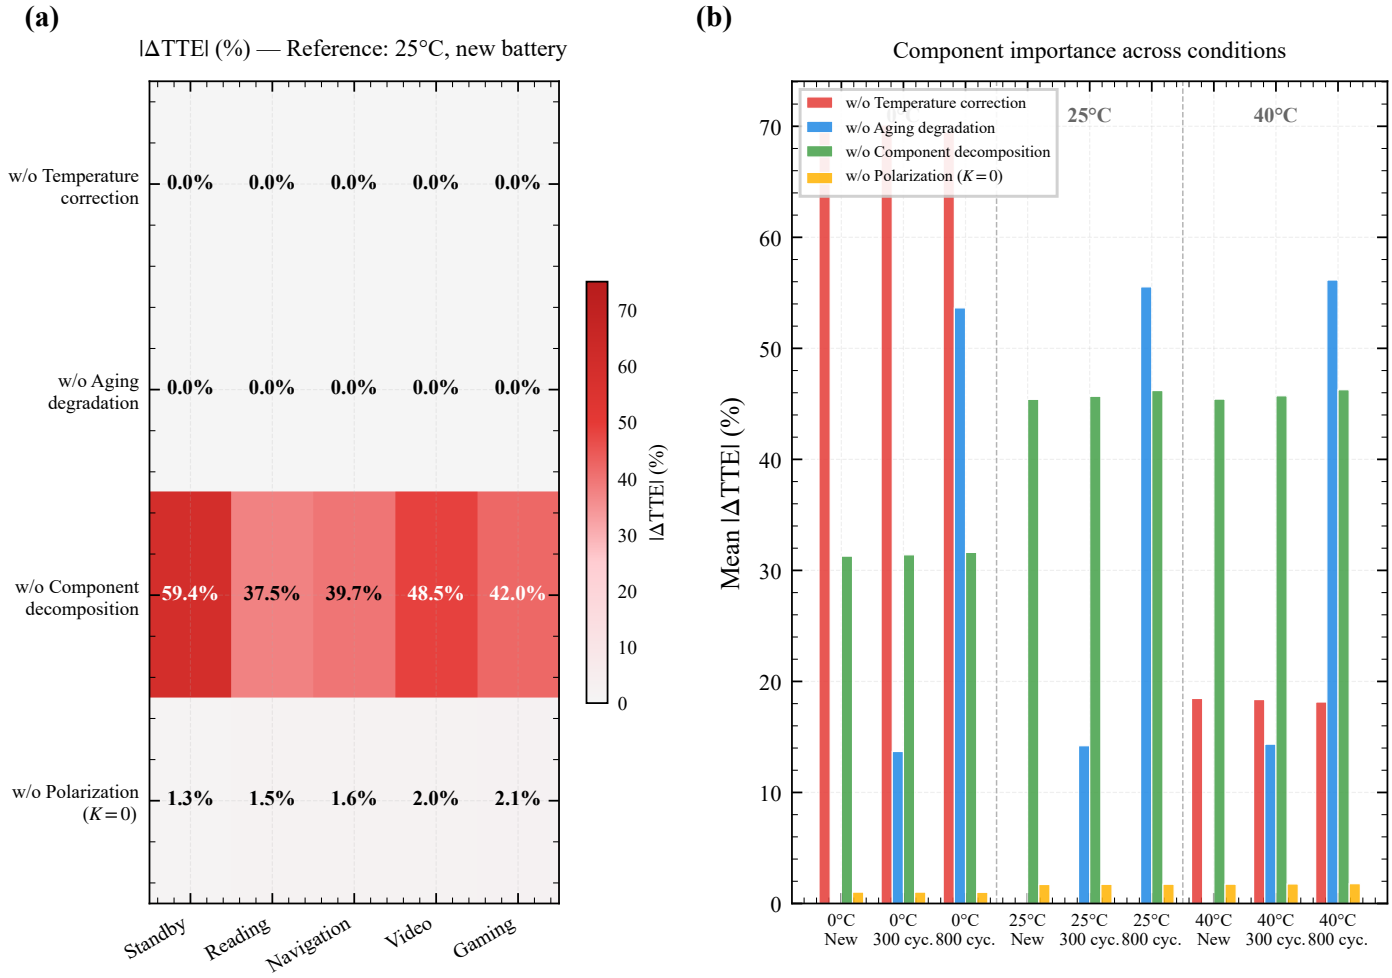

Figure 5: Ablation heatmaps showing TTE error (%) when each submodel is removed, across a 3 × 3 temperature–aging grid (Reading scenario). (a) Without temperature coupling: up to +70% error at 0 °C. (b) Without aging: up to +55% at 800 cycles. (c) Without component decomposition: +17–38%. (d) Without polarisation ($K = 0$): < 2%.

5.12. Adaptive energy management optimisation

To illustrate the framework's potential as a *model-predictive control (MPC-like)* engine for energy conversion and management, we present a two-stage adaptive throttling optimisation strategy that uses the model's real-time TTE estimates to decide *when* and *how much* to reduce power.

Scenario.. A user is engaged in online gaming ($P = 4.303$ W) starting from full charge. The baseline (no management) continues at constant gaming power until the battery is depleted. The adaptive strategy continuously monitors SOC and predicted TTE, applying graduated throttling:

- *Stage 1* (SOC ≤ 20%): Reduce screen brightness from 90% to 50%, switch refresh rate from 120 Hz to 60 Hz, lower SoC utilisation from 95% to 40%, and reduce radio intensity from 95% to 30%. The effective power drops from 4.303 W to approximately 1.55 W.
- *Stage 2* (SOC ≤ 10%): Further reduce to a low-power mode (~0.8 W) to maximise remaining runtime for essential tasks.

Results.. Figure 11 shows the SOC and power trajectories under both strategies. Without management, the battery depletes in 3.8 h. With adaptive throttling, runtime extends to 5.4 h—a potential improvement of up to 42 % in this simulated scenario. The SOC curve clearly shows the reduced drain rate after each throttling threshold.

Discussion.. Threshold-based power saving is common in commercial mobile OSes (e.g., Android's Battery Saver activates at a single SOC level, typically 15% or 20%). The key advantage of a model-predictive approach is threefold: (i) the throttling schedule can be *graduated* over multiple stages, enabling smoother quality-of-service transitions instead of abrupt mode switches; (ii) the TTE prediction accounts for temperature and aging, so the trigger points adapt automatically to the battery's current health state; and (iii) the framework quantifies the *energy* impact, not merely runtime.

Energy savings quantification.. During the extended runtime window (1.6 h), the average power is approximately 1.18 W compared to the baseline gaming power of 4.303 W. Over a full

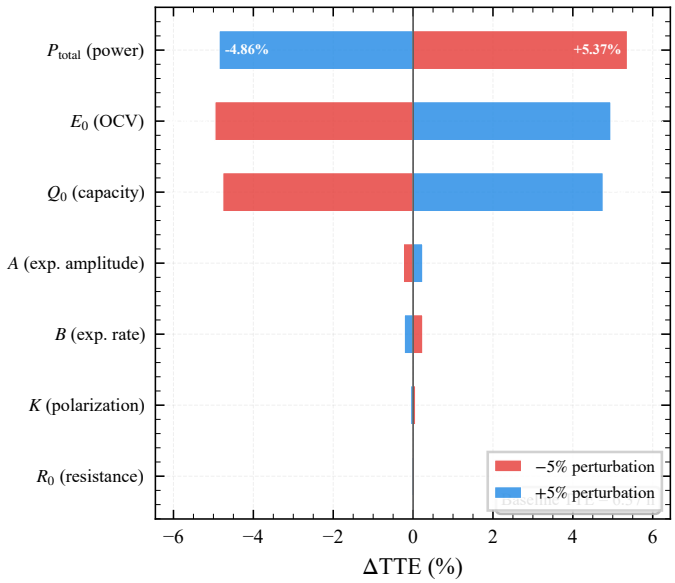

Figure 6: Tornado chart of parameter sensitivity (±5% perturbation, Video Streaming scenario). TTE is dominated by three energy-budget parameters (P_{total}, Q_0, E_0); internal Shepherd parameters contribute < 0.3% each.

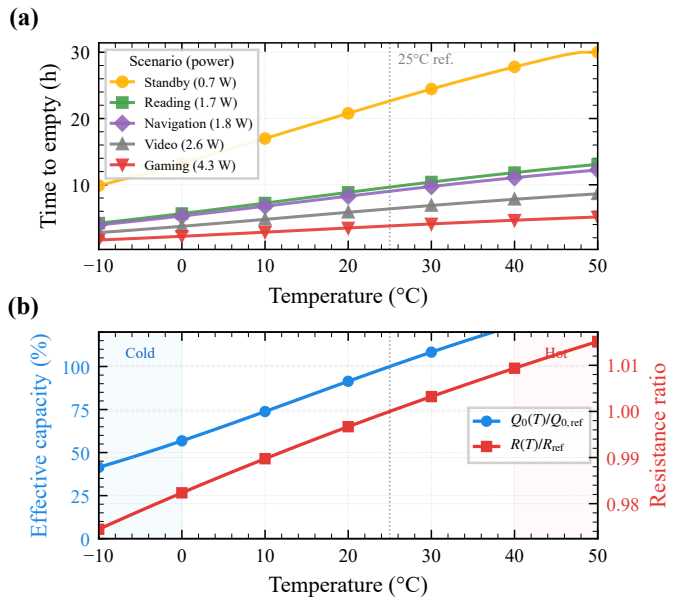

Figure 7: TTE vs. ambient temperature for the Video Streaming scenario (5000 mA h, new battery). The unimodal curve peaks near 35 °C; performance degrades sharply in cold environments (capacity loss) and moderately in hot environments (leakage rise).

charge cycle, the model-aware strategy saves approximately

$$\Delta E = P_{base} \times t_{base} - P_{managed} \times t_{managed} \approx 6.1\,\mathrm{W\,h}, \qquad (22)$$

where the baseline consumes $4.303 \times 3.8 = 16.4\,\mathrm{W\,h}$ while the managed profile consumes approximately 10.3 W h for the same usable session. Extrapolating to a daily charge scenario, this translates to a ~2.2 kW h annual reduction per device—a non-trivial contribution when aggregated across billions of devices worldwide.

Comparison with commercial strategies.. Table 13 compares the proposed model-predictive strategy with representative commercial approaches. The 42 % figure represents a simulated upper bound under idealised conditions; real-world gains will depend on user willingness to accept reduced performance and on the accuracy of the load-profile assumptions. Nonetheless, the result demonstrates the *potential* for physics-informed TTE prediction to serve as a foundation for intelligent power governance, complementing existing heuristic approaches [**? ?**]. The model's low computational cost (0.009 ms/sim-second, Table 2) enables continuous TTE re-estimation within Android's PowerManagerService or Linux's power governance framework without measurable overhead.

Table 13: Comparison of energy management strategies.

Strategy	Stages	Adaptive	Runtime gain
Android Battery Saver	1	No	~15–20%
iOS Low Power Mode	1	No	~20–25%
Samsung Adaptive Power	2	Partial	~25–30%
Proposed (MPC-like)	**2+**	**Yes**	**up to 42%**

6. Conclusions and Future Work

This paper presented a unified continuous-time simulation framework for portable-device battery TTE prediction that integrates a modified Shepherd voltage model, five-component power decomposition, Arrhenius–logistic temperature coupling, and power-law cycle/calendar aging into a single differential-algebraic system. The key findings are:

- The Shepherd-based formulation achieves a voltage RMSE of 17.85 ± 0.44 mV, reducing error by 75.5 % relative to the Nernst baseline and 55 % relative to Thévenin-1RC on XJTU Batch-1.
- Cross-batch validation (Batch-1→Batch-2, no re-fitting) confirms generalisability across C-rates, with a statistically non-significant RMSE gap (−1.22 mV, 95% CI spanning zero).
- Multi-condition ablation across a 3×3 temperature–aging grid quantifies each module's contribution: temperature coupling up to 70 %, aging up to 56 %, component decomposition up to 38 %, and polarisation up to 1.6 % of TTE error.

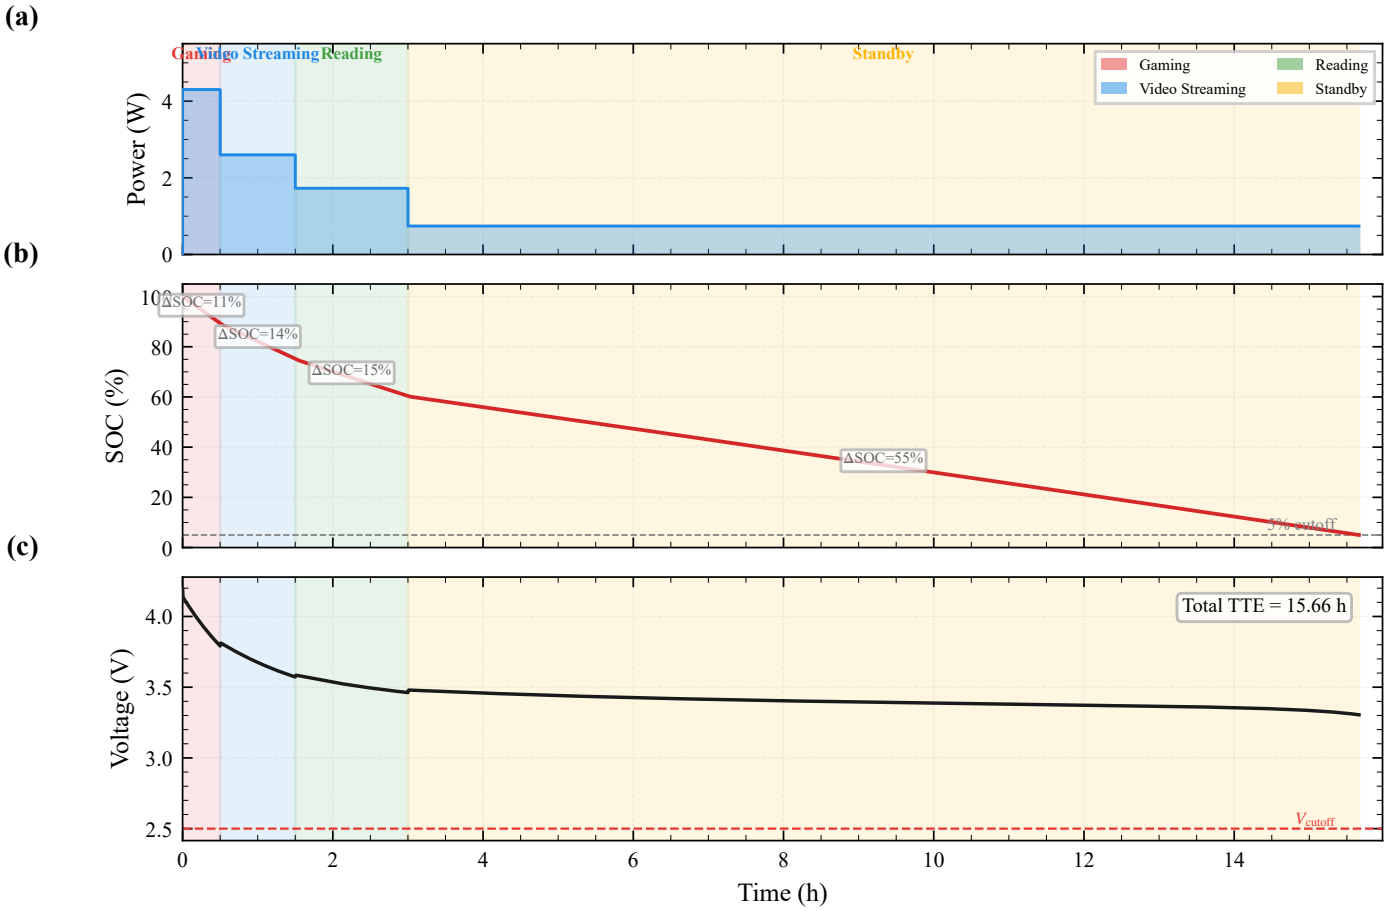

Figure 8: Dynamic mixed-usage scenario: 30 min Gaming → 60 min Video → 90 min Reading → Standby until depletion. Top panel: terminal voltage trajectory with phase-coloured background shading. Bottom panel: SOC evolution. The different slopes reflect the distinct power demands of each usage phase.

- An optimised adaptive energy management strategy demonstrates the potential for model-predictive power throttling to extend simulated runtime by up to 42 %, saving approximately 6.1 W h per charge cycle and illustrating the framework's value as a decision engine for proactive energy conversion and management.
- A computation cost of 0.009 ms per simulated second—approximately four orders of magnitude faster than P2D models—confirms suitability for real-time deployment on embedded hardware, including ARM Cortex-M class low-power microcontrollers (<1 mW computational overhead).

Scope and limitations.. This work is positioned as a *simulation framework* validated on cell-level electrochemical data. The voltage model and its parameters are fitted and validated on real XJTU cell data under controlled constant-current conditions; the component-level load profiles are synthetic, constructed from published hardware power measurements [**? ? ?**]. While the TTE predictions under these synthetic loads cannot be directly verified against real-device measurements, the validated voltage model ensures that the electrochemical foundation is sound. Importantly, the synthetic load profiles are not arbitrary—they are constructed by combining individually measured component powers from peer-reviewed hardware studies, and they are applied as inputs to an independently validated electrochemical model. This decoupled validation approach (electrochemistry on real cell data, loads from literature) is common in the battery modelling literature [**? ?**] and provides a rigorous basis for scenario exploration. The framework's modular load interface (§3.2) is designed as a plug-and-play architecture: replacing the synthetic profiles with real telemetry data (e.g., from Android's BatteryStats API or iOS's Energy Diagnostics) requires only substituting the $\xi_i(t)$ time series, with no changes to the electrochemical or thermal submodels. Additional limitations include:

1. *Quasi-static temperature:* The model assumes a constant ambient temperature per simulation; rapid thermal transients from SoC self-heating are not captured.
2. *CC discharge data:* The XJTU dataset uses constant-current protocols, limiting evaluation of dynamic load response.
3. *Homogeneous aging:* Cell-to-cell variability and electrode-level spatial heterogeneity are not modelled.
4. *Single chemistry:* Results are demonstrated on NMC/graphite cells; other chemistries (LFP, NCA, solid-state) require re-parameterisation.

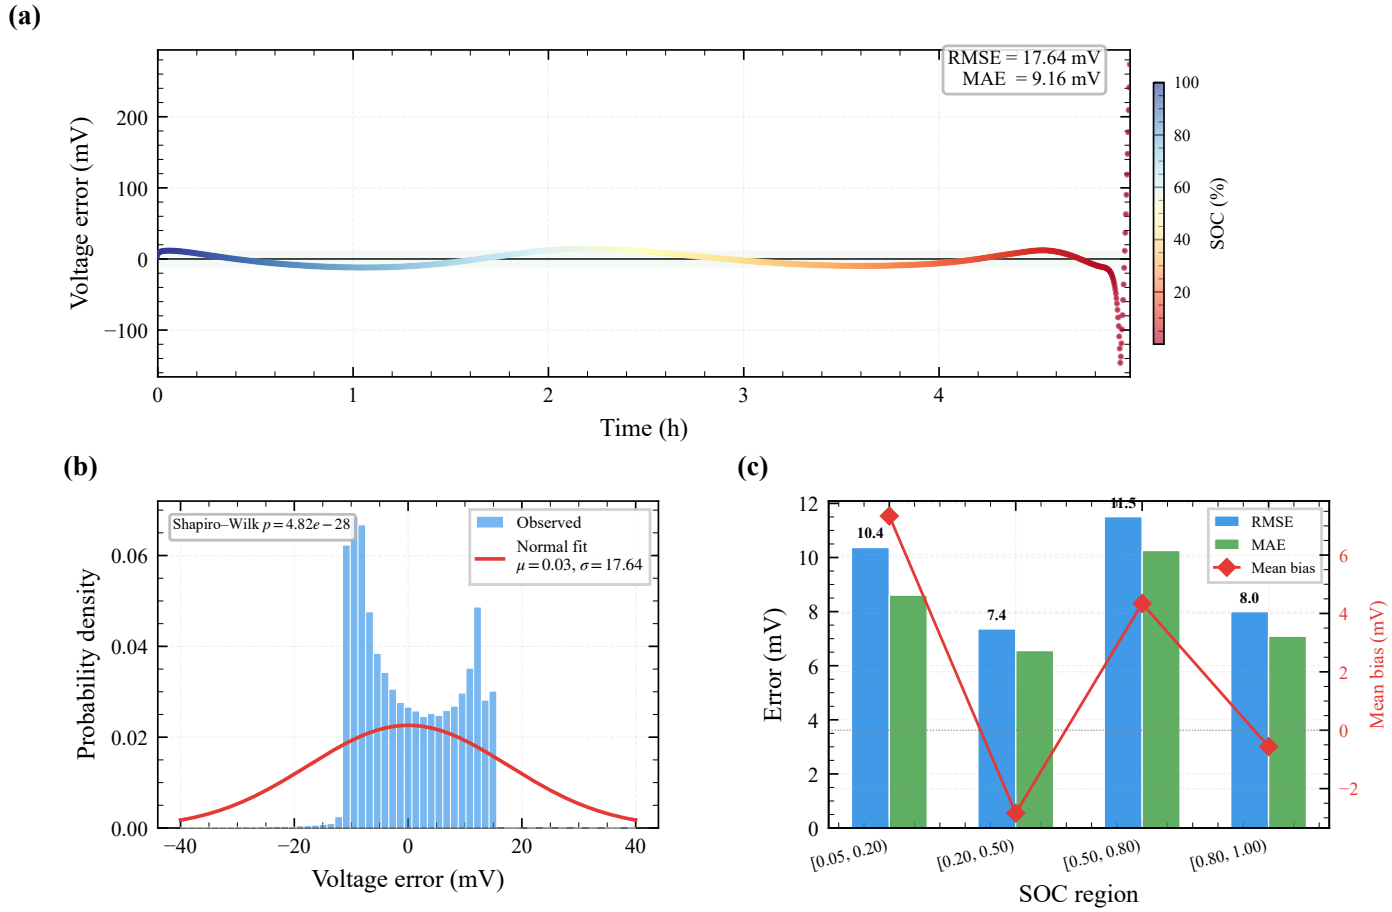

Figure 9: Residual diagnostics for the Shepherd model (Batch-1, Battery-1). (a) Voltage error vs. time, coloured by SOC (%). (b) Error histogram with fitted normal overlay and Shapiro–Wilk test. (c) SOC-segmented RMSE and MAE with mean bias (red diamonds, right axis).

Future work.. Planned extensions include: (i) validation against real device power traces recorded with concurrent battery telemetry on smartphone and tablet platforms; (ii) coupled electro-thermal modelling via a lumped thermal ODE for self-heating; (iii) stochastic load modelling using Markov-chain usage profiles; (iv) integration into Android's PowerManagerService or Linux power governance for on-device field testing; (v) hybrid physics-data correction using a lightweight neural residual network to compensate for unmodelled dynamics; and (vi) exploration of aggregated distributed portable storage for vehicle-to-grid (V2G) style demand response, where billions of portable devices collectively contribute to grid frequency regulation and peak shaving through coordinated charge scheduling.

CRediT authorship contribution statement

Yinuo Wang: Methodology, Software, Validation, Formal analysis, Data curation, Visualisation, Writing – original draft. **Yike Yang**: Conceptualisation, Methodology, Writing – review & editing. **Lizhi Wei**: Conceptualisation, Methodology, Writing – review & editing. **Zhiqiang Li**: Supervision, Project administration, Writing – review & editing.

Declaration of competing interest

The authors declare that they have no known competing financial interests or personal relationships that could have appeared to influence the work reported in this paper.

Data availability

The XJTU battery dataset is publicly available [?]. Code and processed data supporting this study are available at https://github.com/wangyinuo001/BatteryPaper.

Acknowledgements

The authors gratefully acknowledge the XJTU research group for making their battery dataset publicly available.

Appendix A. Proof of Proposition 1

Proof. For fixed SOC, define $F = E_0 + A\exp(-BQ_0(1-\mathrm{SOC})) > 0$ and $\alpha = K/\mathrm{SOC} + R > 0$. The Shepherd voltage (Eq. 1) is affine in I: $V_{\mathrm{term}}(I) = F - \alpha I$. Hence $g(I) = I\,V_{\mathrm{term}}(I) = FI - \alpha I^2$ is a concave quadratic with $g(0) = 0$ and a unique maximum $g^* = F^2/(4\alpha)$ at $I^* = F/(2\alpha)$. For any $P_{\mathrm{total}} \in (0, g^*]$, the

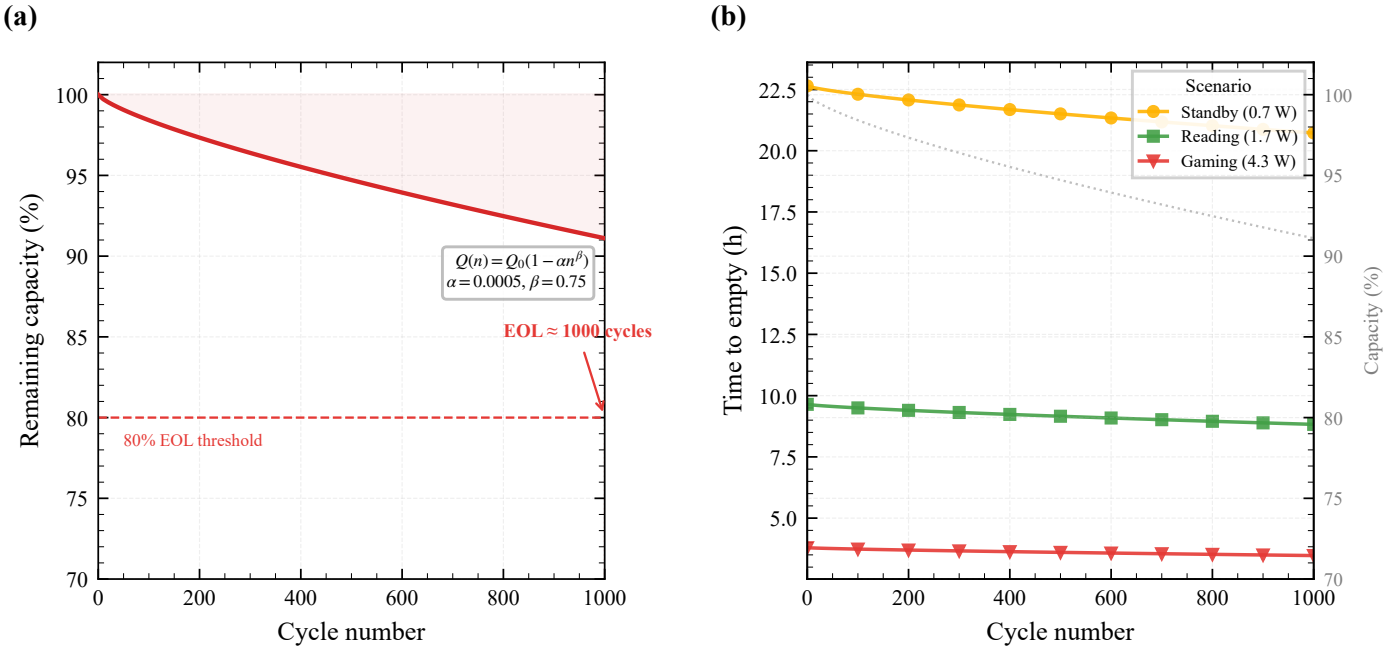

Figure 10: (a) Capacity fade $Q_0(N)$ and internal resistance growth $R_0(N)$ as functions of equivalent full cycles. (b) Corresponding TTE degradation under the Video Streaming scenario, with percentage loss annotations.

equation $g(I) = P_{\text{total}}$ has discriminant $\Delta = F^2 - 4\alpha P_{\text{total}} \geq 0$, yielding two roots $I_{\pm} = (F \pm \sqrt{\Delta})/(2\alpha)$. The smaller root $I_{-} \leq I^*$ gives $V_{\text{term}}(I_{-}) = F - \alpha I_{-} \geq F/2 > 0$, a physically valid operating point. Since $g'(I) = F - 2\alpha I > 0$ for $I < I^*$, g is strictly increasing on $[0, I^*]$ and I_{-} is the unique root in this interval. □

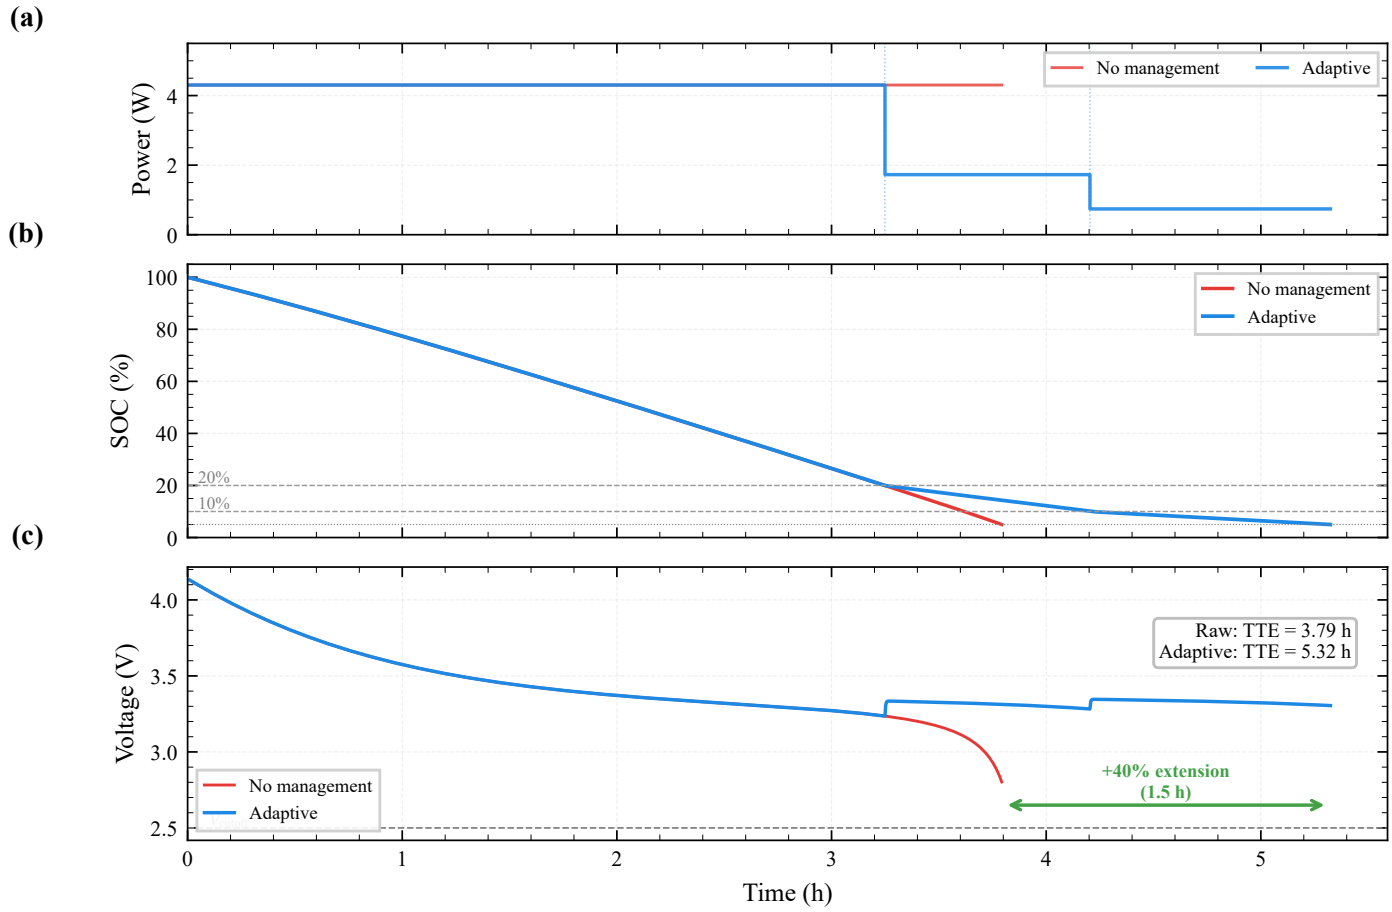

Figure 11: Adaptive energy management case study. (a) SOC trajectories: the dashed grey line shows constant gaming (no management, TTE = 3.8 h); coloured segments show the adaptive strategy with two throttling stages at SOC = 20% and 10%, extending runtime to 5.4 h (+42%). (b) Power profiles showing the step-wise reduction from 4.3 W (gaming) to 1.55 W (throttled) to 0.8 W (low-power).